

# Grounded Language-Image Pre-training

Liunian Harold Li, Pengchuan Zhang, Haotian Zhang, Jianwei Yang, Chunyuan Li, Yiwu Zhong, Lijuan Wang, Lu Yuan, Lei Zhang, Jenq-Neng Hwang, Kai-Wei Chang, Jianfeng Gao

**CVPR 2022**

Presenter: Yi-Ting Li (李易庭)

Date: Feb. 15, 2023



*Intelligent Knowledge Management Laboratory  
Department of Computer Science and Information Engineering  
National Cheng Kung University, Tainan, Taiwan, R.O.C*

# Outline

## Task Definition

1. Object Detection
2. Phrase Grounding

## Related Work

1. DERT - Facebook
2. MDETR - Facebook

## Methodology

1. Reformulating object detection as phrase grounding
2. Language-Aware Deep Fusion
3. Pre-training with Scalable Semantic-Rich Data

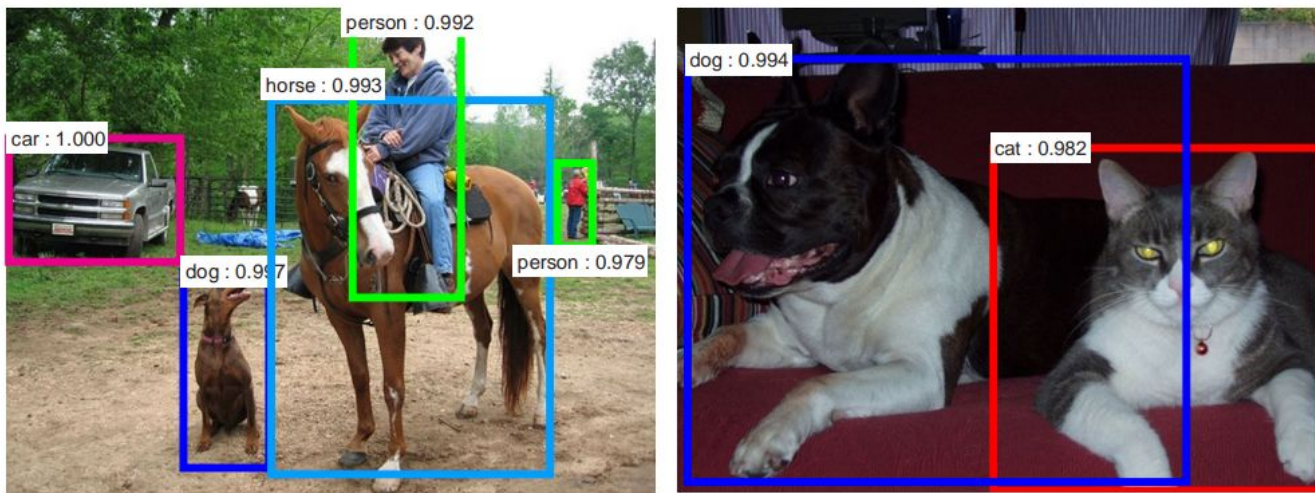
## Transfer learning result

1. Zero-shot
2. Prompt tuning
3. Linear probing

## SWOT


# Task Definition - Object Detection

Object detection is the task of detecting instances of **objects** of a **certain class** within an image.



# Task Definition - Phrase Grounding

Given an **image** and a **corresponding caption**, the Phrase Grounding task aims to ground each **entity** mentioned by a noun phrase in the caption to a region in the image.



A tattooed woman with a green dress and yellow backpack holding a water bottle is walking across the street.


# Task Definition - Problem and Solution

## Problem

1. While training an object detection model, the model will be given an **image** and **bounding box** with **text label**. Therefore the annotation fee will be costly.
2. The object detection model can only detect instances of objects of a particular class, so for the labels **not in** the training dataset, the model can't figure them out.

## Solution - Treating the object detection tasks as phrase grounding


1. We could easily collect image-caption pairs from the Internet.
2. By giving prompts(caption) to control the detection target.



Prompt : person. bicycle.  
car. motorcycle...



Prompt : aerosol can...  
lollipop... pendulum...




Prompt : raccoon



Prompt : pistol



Prompt : there are some  
holes on the road



Prompt : person. dog.

# Task Definition - Image Caption Pair

## In GLIP:

Like CLIP, unknown

## Current model:

Use BLIP



$T_w$ : "from bridge near my house"

$T_s$ : "a flock of birds flying over a lake at sunset"




$T_w$ : "in front of a house door in Reichenfels, Austria"


$T_s$ : "a potted plant sitting on top of a pile of rocks"



"blue sky bakery in sunset park"



"chocolate cake with cream frosting and chocolate sprinkles on top"



# Related Work

# DETR

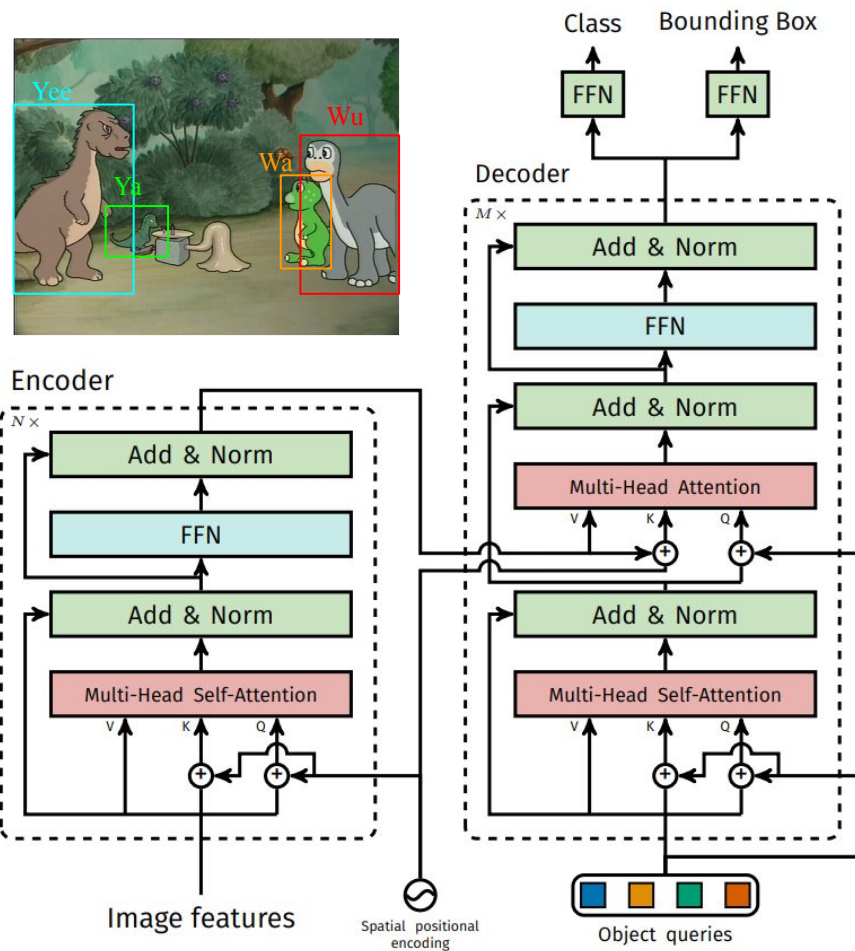
## Traditional Object Detection Problem

Need so many **hand-design** component:

1. Proposals(2-stage, e.g., R-CNN family)
2. Anchor or Window centers(1-stage, e.g., Yolo family)
3. NMS (Non-maximum suppression)

## Contribution

1. **Without** hand-design components anymore.
2. The **object detection set(cls, box) prediction loss**
3. Parallel Decoding (Transformer)





# DETR

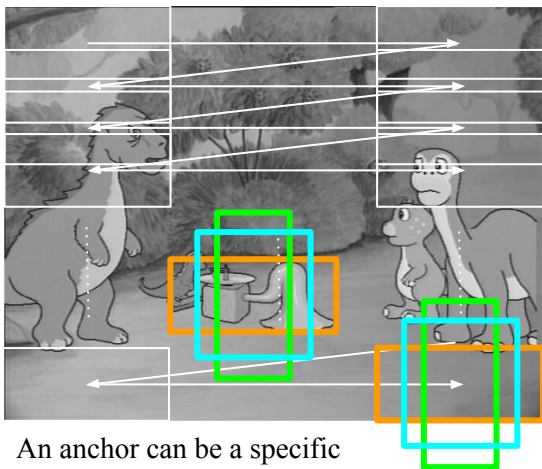
## DETR implement with PyTorch

```
1  import torch
2  from torch import nn
3  from torchvision.models import resnet50
4
5  class DETR(nn.Module):
6
7      def __init__(self, num_classes, hidden_dim, nheads,
8                  num_encoder_layers, num_decoder_layers):
9          super().__init__()
10         # We take only convolutional layers from ResNet-50 model
11         self.backbone = nn.Sequential(*list(resnet50(pretrained=True).children())[:-2])
12         self.conv = nn.Conv2d(2048, hidden_dim, 1)
13         self.transformer = nn.Transformer(hidden_dim, nheads,
14                                           num_encoder_layers, num_decoder_layers)
15         self.linear_class = nn.Linear(hidden_dim, num_classes + 1)
16         self.linear_bbox = nn.Linear(hidden_dim, 4)
17         self.query_pos = nn.Parameter(torch.rand(100, hidden_dim))
18         self.row_embed = nn.Parameter(torch.rand(50, hidden_dim // 2))
19         self.col_embed = nn.Parameter(torch.rand(50, hidden_dim // 2))
20
21     def forward(self, inputs):
22         x = self.backbone(inputs)
23         h = self.conv(x)
24         H, W = h.shape[-2:]
25         pos = torch.cat([
26             self.col_embed[:W].unsqueeze(0).repeat(H, 1, 1),
27             self.row_embed[:H].unsqueeze(1).repeat(1, W, 1),
28         ], dim=-1).flatten(0, 1).unsqueeze(1)
29         h = self.transformer(pos + h.flatten(2).permute(2, 0, 1),
30                             self.query_pos.unsqueeze(1))
31         return self.linear_class(h), self.linear_bbox(h).sigmoid()
```

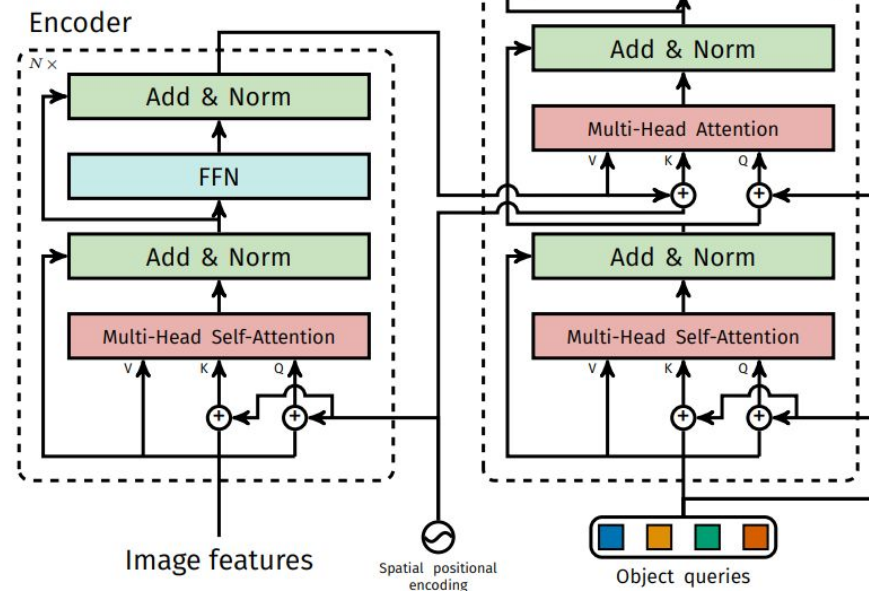

# DETR

Aiden: What's the **meaning** of Object Queries?

Answer: Like a set of anchors in traditional object detection models but learnable.




An anchor can be a specific width and height rectangle



# DETR

Each box **prediction** is represented as a **point** with the coordinates of its center in the 1-by-1 square normalized by each image size.

1. Green: small bboxes
2. Red: large horizontal bboxes
3. Blue: large vertical bboxes



# DETR

Question: How to compute the loss? The number of output pairs always **larger** than the number of ground truth.

Answer: Optimal **bipartite matching**

Step1) Construct the “Cost Matrix” by match loss function

$$-\mathbb{1}_{\{c_i \neq \emptyset\}} \hat{p}_{\sigma(i)}(c_i) + \mathbb{1}_{\{c_i \neq \emptyset\}} \mathcal{L}_{\text{box}}(b_i, \hat{b}_{\sigma(i)})$$

Step2) Apply **Hungarian Algorithm** to find the perfect one-to-one matching

Step3) Compute the loss value

$$\sum_{i=1}^N \left[ -\log \hat{p}_{\hat{\sigma}(i)}(c_i) + \mathbb{1}_{\{c_i \neq \emptyset\}} \mathcal{L}_{\text{box}}(b_i, \hat{b}_{\hat{\sigma}(i)}) \right]$$


		Ground Truth			
		L <sub>1</sub>	L <sub>2</sub>	L <sub>3</sub>	L <sub>4</sub>
Predict	O <sub>1</sub>	1	2	3	4
	O <sub>2</sub>	2	2	4	1
	O <sub>3</sub>	3	2	1	4
	O <sub>4</sub>	4	4	2	3
	O <sub>5</sub>	3	1	2	2

Cost Matrix




# DETR

The encoder self-attention



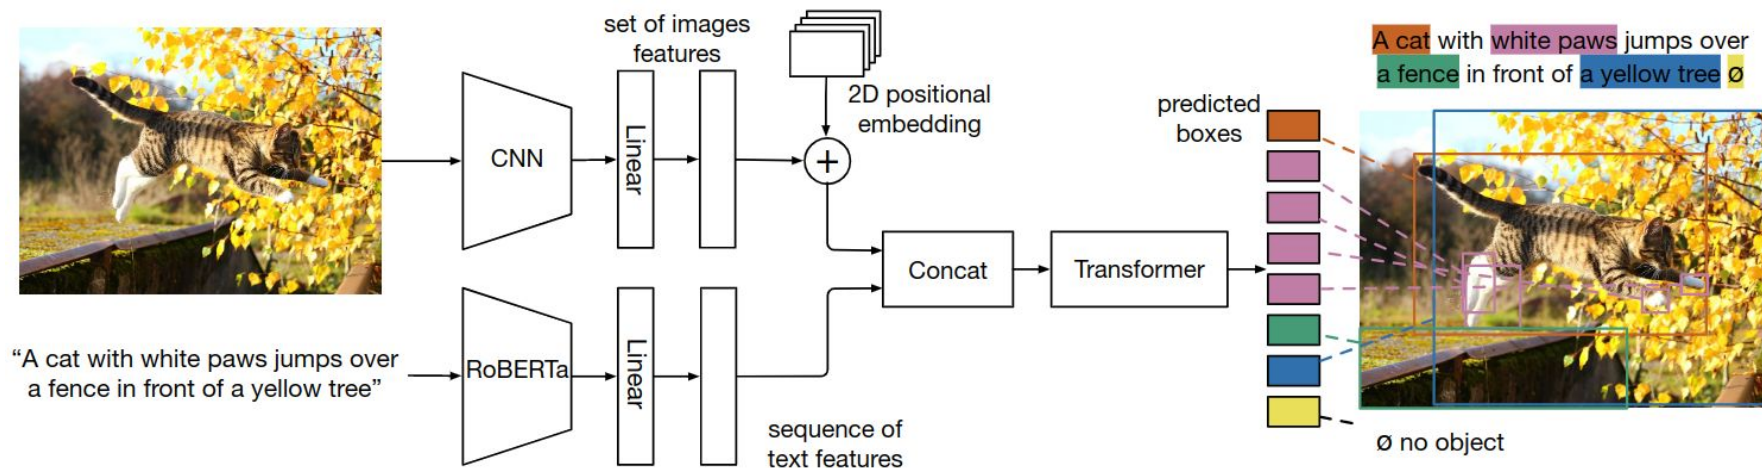
The decoder attention



# MDETR

## 英嘉學長: Why CNN is needed? Why don't just use Transformer?

Answer: Reason for input. The input for Transformer will be too long, if we flatten the image as the input directly; the pre-trained CNNs(like ResNet) pay more attention to the local feature(like contour)



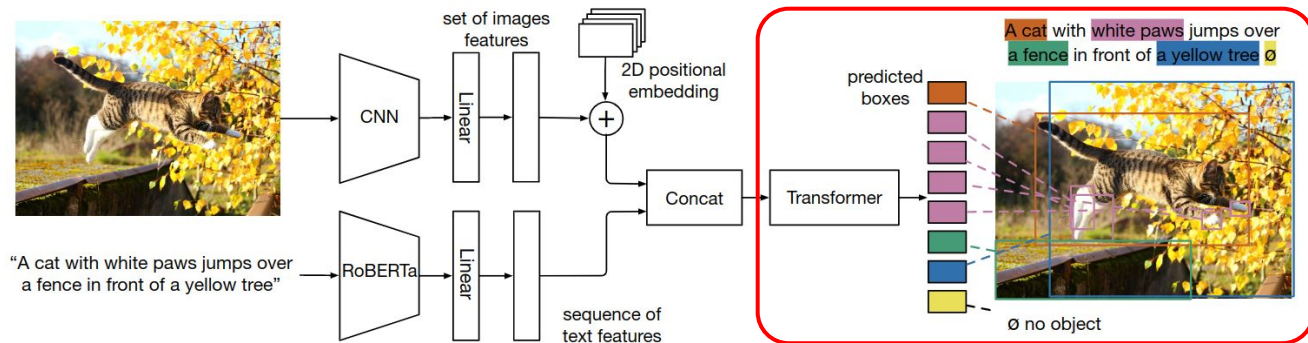
[9] Kamath, Aishwarya, et al. "Mdetr-modulated detection for end-to-end multi-modal understanding." ICCV. 2021.

[10] Dosovitskiy, Alexey, et al. "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale." ICLR. 2020.

# MDETR

## Soft token prediction

1. RoBERTa
2. The model is trained to predict a **uniform distribution** over all token positions that correspond to the object.




**Aiden:** How to design the loss function? Answer: as a Multi-label classification (by CE).

## Contrastive alignment(InfoNCE)

Positive Pair:  ,  A cat

Negative Pair:  , with white ...  $\emptyset$

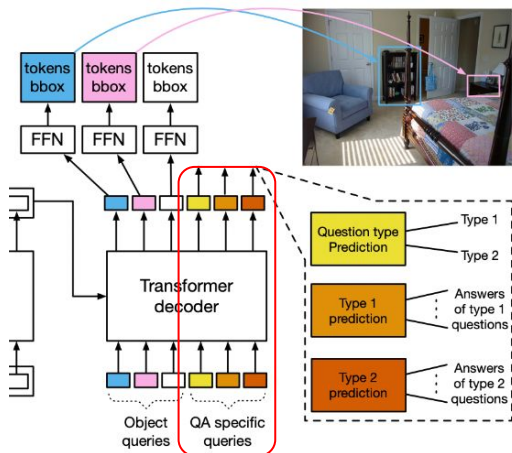




# MDETR

## Extend model for VQA (GQA dataset)

Provide QA specific queries REL, OBJ, GLOBAL, CAT and ATTR in addition to the object queries as input to the decoder



Question: "What is on the table?"

Answer:

- Question type prediction: **OBJ**
- Type 1 prediction: **Laptop**

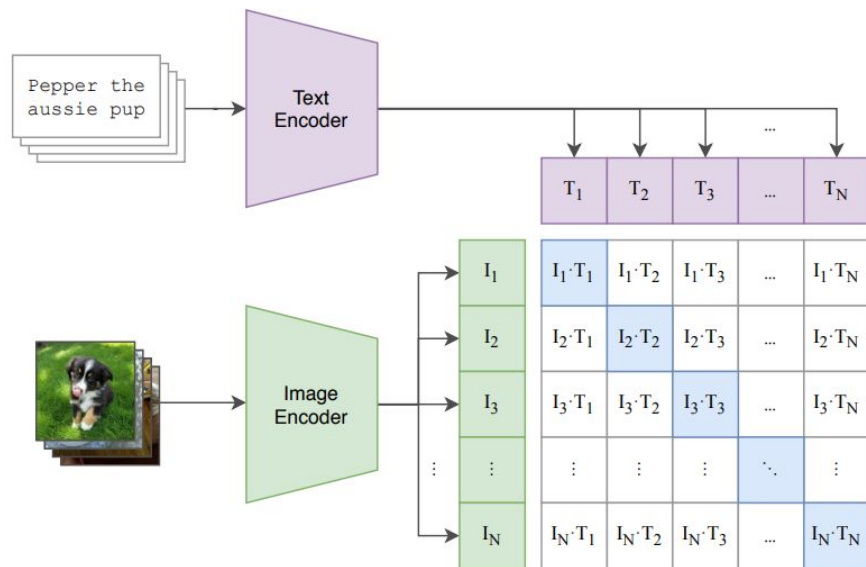


# MDETR

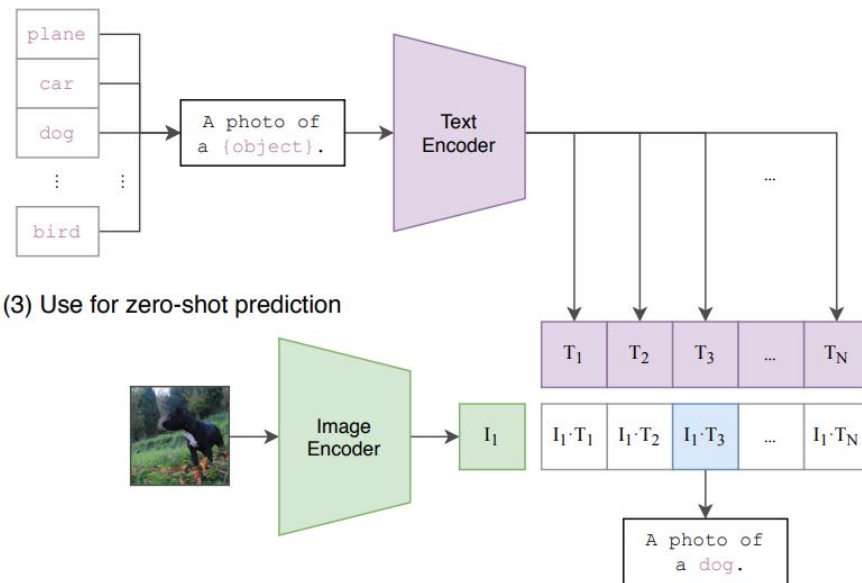
Type	Open/Binary	Semantic	Structural	Form	Example
queryGlobal	open	query	<u>global</u>	<u>select: scene/query: type</u>	<u>How is the weather in the image?</u>
verifyGlobal	binary	verify	global	select: scene/verify type: attr	Is it cloudy today?
chooseGlobal	open	query	global	select: scene/choose type: a b	Is it sunny or cloudy?
queryAttr	open	query	<u>attribute</u>	<u>select: obj/.../query: type</u>	<u>What color is the apple?</u>
verifyAttr	binary	verify	attribute	select: obj/.../verify type: attr	Is the apple red?
verifyAttr	binary	logical	attribute	select: obj/.../verify t1: a1/verify t2: a2/and	Is the apple red and shiny?
chooseAttr	open	choose	attribute	select: obj/.../choose type: a b	Is the apple green or red?
exist	binary	verify	<u>object</u>	<u>select: obj/.../exist</u>	<u>Is there an apple in the picture?</u>
existRel	binary	verify	<u>relation</u>	<u>select: subj/.../relate (rel): obj/exist</u>	<u>Is there an apple on the black table?</u>
logicOr	binary	logical	object	select: obj1/.../exist/select: obj2/.../exist/or	Do you see either an apple or a banana there?
logicAnd	binary	logical	obj/attr	select: obj1/.../exist/select: obj2/.../exist/and	Do you see both green apples and bananas there?
queryObject	open	query	<u>category</u>	<u>select: category/.../query: name</u>	<u>What kind of fruit is on the table?</u>
chooseObject	open	choose	category	select: category/.../choose: a b	What kind of fruit is it, an apple or a banana?

# Related Work - CLIP


## (1) Contrastive pre-training



## (2) Create dataset classifier from label text



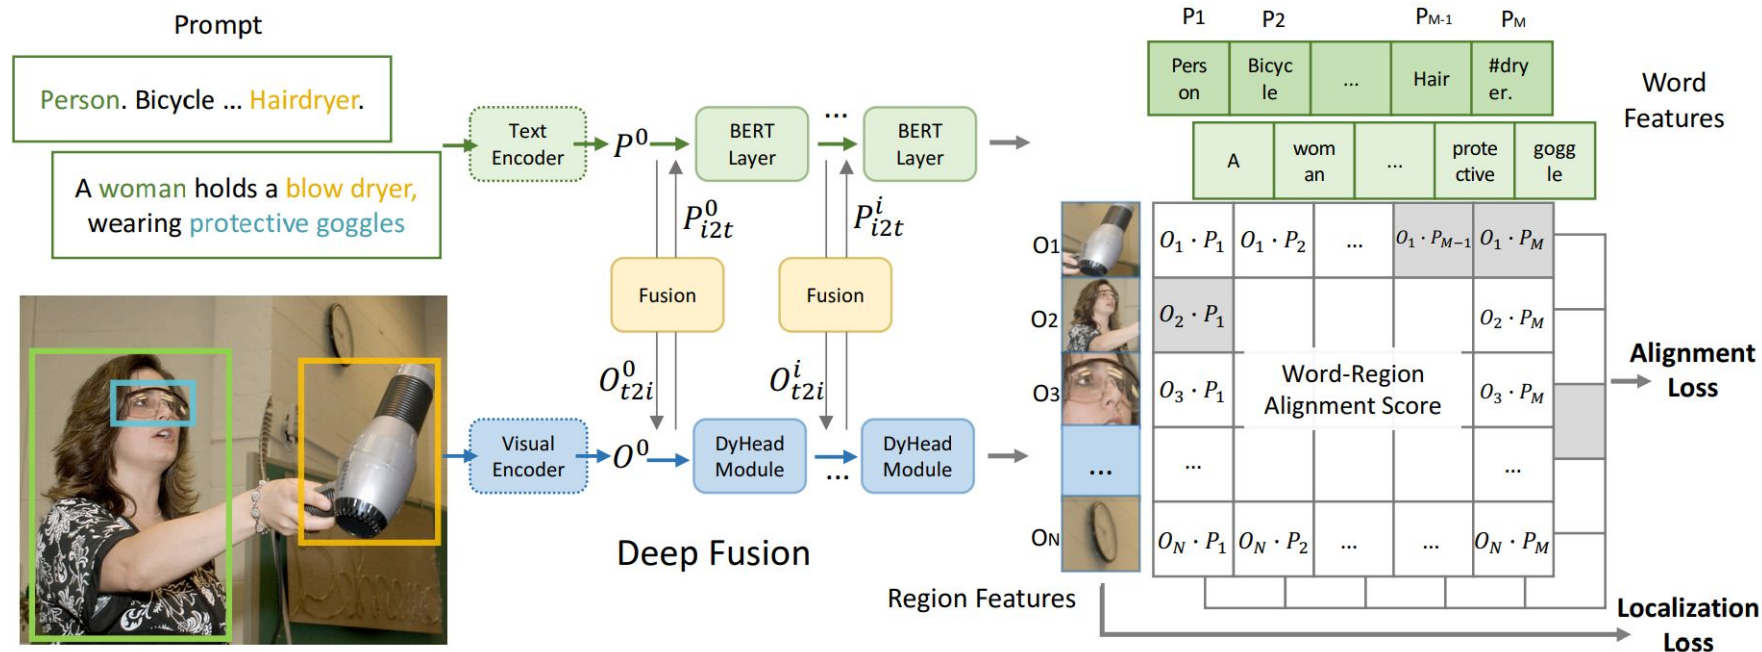
## (3) Use for zero-shot prediction



# Methodology

# Reformulating object detection as phrase grounding

## Object Detection as Phrase Grounding



# Reformulating object detection as phrase grounding

**From Object Detection**  $\mathcal{L} = \mathcal{L}_{\text{cls}} + \mathcal{L}_{\text{loc}}$

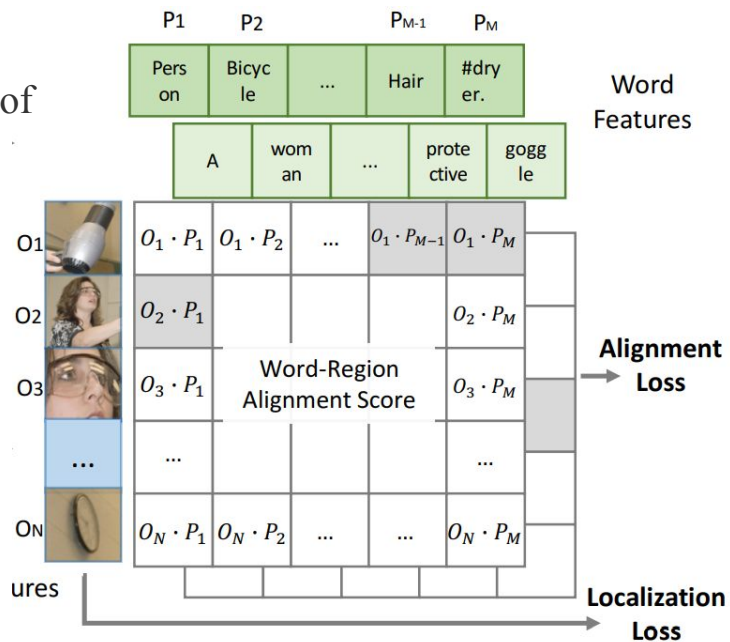
**To Phrase Grounding**  $\mathcal{L} = \mathcal{S}_{\text{ground}} + \mathcal{L}_{\text{loc}}$

**Problem:** the number of logits will always **larger** than the number of phrases in the text prompt due to following reasons:

1. some phrases contain multiple words, e.g. “traffic light”.
2. some are the added tokens, e.g., “Detect:”, “,”.
3. some single-word phrases are splitted into multiple (sub)-word tokens, e.g., “toothbrush” to “tooth#” and “#brush”.
4. [NoObj] token is added at the end of the tokenized sequence.

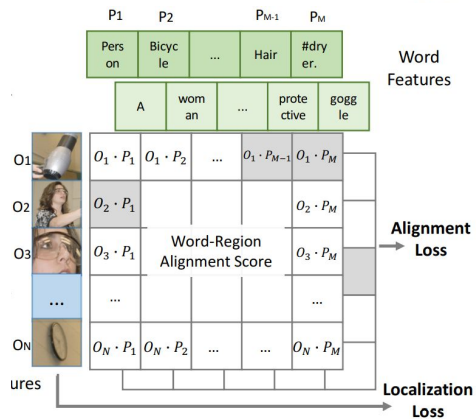
## Solution:

If one phrase is positive match make **all** sub-words positive match.  
During inference, **average** token probabilities as the phrase probability.



# Reformulating object detection as phrase grounding

## Object Detection $\mathcal{L} = \mathcal{L}_{\text{cls}} + \mathcal{L}_{\text{loc}}$



## To Phrase Grounding $\mathcal{L} = \mathcal{S}_{\text{ground}} + \mathcal{L}_{\text{loc}}$

1. some phrases contain multiple words, e.g. “traffic light”
2. some are the added tokens, e.g., “Detect:”, “,”
3. some single-word phrases are splitted into multiple (sub)-word tokens, e.g., “toothbrush” to “tooth#” and “#brush”
4. [NoObj] token is added at the end of the tokenized sequence.

$$O = \text{Enc}_I(\text{Img}) \quad (1)$$

$$S_{\text{cls}} = OW^T \quad (2)$$

$$\mathcal{L}_{\text{cls}} = \text{loss}(S_{\text{cls}}; T) \quad (3)$$

$O \in \mathbb{R}^{N \times d}$  = object/region/box features of the input image

$W \in \mathbb{R}^{c \times d}$  = weight matrix of the box classifier

$S_{\text{cls}} \in \mathbb{R}^{N \times c}$  = logits of the classification

$T \in \{0, 1\}^{N \times c}$  = target matching between regions and classes

$\text{loss}(S; T)$  = cross-entropy loss for two-stage detectors

$c$  = number of classes

$$O = \text{Enc}_I(\text{Img}) \quad (4)$$

$$P = \text{Enc}_L(\text{Prompt}) \quad (5)$$

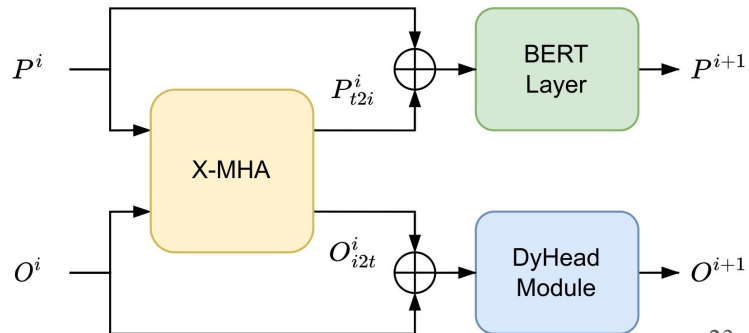
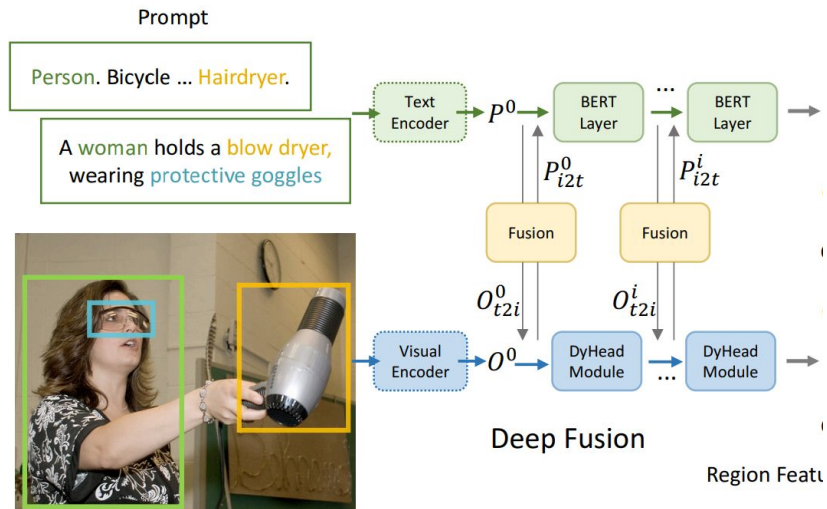
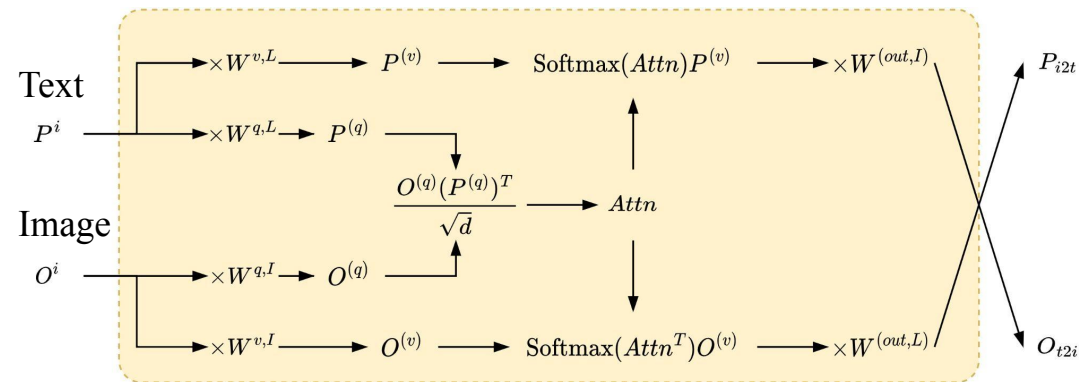
$$S_{\text{ground}} = OP^T \quad (6)$$

$P \in \mathbb{R}^{M \times d}$  = word/token features from language encoder

# Language-Aware Deep Fusion

Simply fuse Bert-Layer with DyHead-Layer by  
Cross-Modal Multi-head Attention (X-MHA)

$$\begin{aligned}
 O^{(q)} &= OW^{(q,I)}, P^{(q)} = PW^{(q,L)}, Attn = O^{(q)}(P^{(q)})^\top / \sqrt{d}, \\
 P^{(v)} &= PW^{(v,L)}, O_{t2i} = \text{SoftMax}(Attn)P^{(v)}W^{(out,I)}, \\
 O^{(v)} &= OW^{(v,I)}, P_{i2t} = \text{SoftMax}(Attn^\top)O^{(v)}W^{(out,L)},
 \end{aligned}$$





# Pre-training with Scalable Semantic-Rich Data


1. Combine all the object detection dataset to train a GLIP model as the teacher model.
2. Scrape the image-text pair from the Internet as augmentation.
3. The inference output from teacher model as pseudo-labels to train the student model.

Model	Backbone	Pre-Train Data	Zero-Shot 2017val	Fine-Tune 2017val / test-dev
Faster RCNN	RN50-FPN	-	-	40.2 / -
Faster RCNN	RN101-FPN	-	-	42.0 / -
DyHead-T [10]	Swin-T	-	-	49.7 / -
DyHead-L [10]	Swin-L	-	-	58.4 / 58.7
DyHead-L [10]	Swin-L	O365,ImageNet21K	-	60.3 / 60.6
SoftTeacher [65]	Swin-L	O365,SS-COCO	-	60.7 / 61.3
DyHead-T	Swin-T	O365	43.6	53.3 / -
w/o deep fusion GLIP-T (A)	Swin-T	O365	42.9	52.9 / -
w/ deep fusion GLIP-T (B)	Swin-T	O365	44.9	53.8 / -
Teacher model GLIP-T (C)	Swin-T	O365,GoldG	<b>46.7</b>	55.1 / -
GLIP-T	Swin-T	O365,GoldG,Cap4M	46.3	54.9 / -
GLIP-T	Swin-T	O365,GoldG,CC3M,SBU	46.6	<b>55.2</b> / -
GLIP-L	Swin-L	FourODs,GoldG,Cap24M	<b>49.8</b>	<b>60.8</b> / 61.0
GLIP-L	Swin-L	FourODs,GoldG+,COCO	-	- / <b>61.5</b>




# Prompt Tuning


By only tuning the prompt embeddings, GLIP-T and GLIP-L can achieve performance close to full-model tuning, allowing for efficient deployment.



Prompt: ... **stingray** ...



Prompt: ... **stingray**, which is flat and round...



# SWOT

## Strengths

1. Deployment efficiency
2. Without hand-designed components to detect objects.
3. Fusion BERT with DyHead to get better cross-modal features.

## Opportunities

1. We can design a prompt for VQA to get the answer bounding box.

## Weaknesses

1. Memory costly and inference slowly

Model	Fusion	Inference (P100)		Train (V100)	
		Speed	Memory	Speed	Memory
GLIP-T	✗	4.84 FPS	1.0 GB	2.79 FPS	11.5 GB
	✓	2.52 FPS	2.4 GB	1.62 FPS	16.0 GB
GLIP-L	✗	0.54 FPS	4.8 GB	1.27 FPS	19.7 GB
	✓	0.32 FPS	7.7 GB	0.88 FPS	23.4 GB

## Threats

1. The model may can not understand the specific prompt pattern like: <text\_1>